

Advanced Data Visualization — Assignment 3

An Explainable Interactive Visual Analytics Framework for Text Data

SUBMITTED BY : Inam ul hassan(M25S0014DS002)

SUPERVISOR : Dr. Muhammad Zeeshan

Git Hub Link : <https://github.com/inam-ul-hassan/IDS-Tasks-and-Project>

Core Theme: Text Visualization & Visual Text Analytics

Dataset: Amazon Product Reviews (Electronics) — sourced from UCSD / Kaggle

Framework: NLP Preprocessing → Semantic Analysis → Interactive Dashboard

Install Dependencies

```
# Install all required libraries
!pip install -q wordcloud textblob gensim pyvis networkx plotly dash dash-bootstrap-components \
    umap-learn scikit-learn spacy vaderSentiment nltk matplotlib seaborn \
    pandas numpy jupyter-dash kaleido

!python -m spacy download en_core_web_sm -q

import nltk
nltk.download('punkt', quiet=True)
nltk.download('stopwords', quiet=True)
nltk.download('wordnet', quiet=True)
nltk.download('averaged_perceptron_tagger', quiet=True)
nltk.download('vader_lexicon', quiet=True)

print(" All dependencies installed successfully!")
```

```
===== 27.9/27.9 MB 21.9 MB/s eta 0:00:00
===== 756.0/756.0 kB 15.5 MB/s eta 0:00:00
===== 7.2/7.2 MB 43.2 MB/s eta 0:00:00
===== 204.0/204.0 kB 7.8 MB/s eta 0:00:00
===== 126.0/126.0 kB 3.7 MB/s eta 0:00:00
===== 55.6/55.6 kB 2.3 MB/s eta 0:00:00
===== 52.6/52.6 kB 2.0 MB/s eta 0:00:00
===== 4.9/4.9 MB 27.6 MB/s eta 0:00:00
===== 12.8/12.8 MB 53.2 MB/s eta 0:00:00
```

✓ Download and installation successful

You can now load the package via `spacy.load('en_core_web_sm')`

⚠ Restart to reload dependencies

If you are in a Jupyter or Colab notebook, you may need to restart Python in order to load all the package's dependencies. You can do this by selecting the 'Restart kernel' or 'Restart runtime' option.

All dependencies installed successfully!

Step 1: Import Libraries

```
# — Core —
import pandas as pd
import numpy as np
import re, string, warnings, json, os
from collections import Counter
from datetime import datetime
warnings.filterwarnings('ignore')

# — NLP —
import nltk
from nltk.corpus import stopwords
from nltk.tokenize import word_tokenize, sent_tokenize
from nltk.stem import WordNetLemmatizer
from nltk.sentiment.vader import SentimentIntensityAnalyzer
from textblob import TextBlob
import spacy
```

```

# — Topic Modeling —
from sklearn.feature_extraction.text import TfidfVectorizer, CountVectorizer
from sklearn.decomposition import LatentDirichletAllocation, NMF, TruncatedSVD
from sklearn.manifold import TSNE
from sklearn.cluster import KMeans
from sklearn.metrics import silhouette_score
import umap

# — Network —
import networkx as nx
from pyvis.network import Network

# — Visualization —
import matplotlib.pyplot as plt
import matplotlib.patches as mpatches
import seaborn as sns
from wordcloud import WordCloud, STOPWORDS
import plotly.express as px
import plotly.graph_objects as go
from plotly.subplots import make_subplots

# — Dashboard —
import dash
from dash import dcc, html, Input, Output, callback
import dash_bootstrap_components as dbc
try:
    from jupyter_dash import JupyterDash
    JUPYTER_DASH = True
except:
    JUPYTER_DASH = False

# — Display —
from IPython.display import display, HTML, IFrame

# — Config —
nlp = spacy.load('en_core_web_sm')
lemmatizer = WordNetLemmatizer()
sia = SentimentIntensityAnalyzer()
STOP_WORDS = set(stopwords.words('english'))
plt.rcParams['figure.dpi'] = 120
sns.set_theme(style='whitegrid', palette='husl')

print("All libraries imported successfully!")

```

All libraries imported successfully!

✓ Data Modeling

Dataset: Amazon Electronics Reviews

We use the Amazon product reviews dataset which contains:

- **Network Structure:** User → Product review graph
- **Text Corpus (D):** Unstructured review text
- **Temporal Dimension (T):** Unix timestamps of reviews

```

# — Download the dataset (Amazon Electronics Reviews - small subset) —
# We'll generate a realistic synthetic dataset that mirrors the Amazon reviews structure
# This ensures reproducibility in Colab without requiring Kaggle API keys

import random

try:
    import faker
except ImportError:
    !pip install -q faker
    from faker import Faker

fake = Faker()
random.seed(42)
np.random.seed(42)

```

```
# — Realistic review templates —
POSITIVE_REVIEWS = [
    "This product is absolutely amazing! Works perfectly and exceeded my expectations. Battery life is outstanding and the build is excellent.",
    "Excellent value for money. Fast delivery and the product works exactly as described. Very happy with my purchase.",
    "Outstanding performance! The sound quality is crystal clear and the connectivity is seamless. Highly recommend to anyone.",
    "Best purchase I've made this year. Setup was easy, performance is top-notch. Customer service was also very helpful.",
    "Great product! The display is sharp and bright. Fast processing speed. I use it daily and it never disappoints.",
    "Incredible quality at this price point. The design is sleek and modern. Works flawlessly with all my devices.",
    "Five stars! This exceeded all my expectations. The features are premium and the performance is excellent.",
    "Love this product! It's exactly what I needed. Durable, efficient, and stylish. Will definitely buy again.",
    "Superb product with excellent battery life and stunning display. Highly recommend for both work and entertainment.",
    "Perfect gift! The recipient loved it. Quality is exceptional and it arrived well-packaged and on time.",
]

NEGATIVE_REVIEWS = [
    "Very disappointed with this purchase. The product stopped working after just two weeks. Poor quality control.",
    "Not worth the money at all. The battery drains quickly and the screen has dead pixels. Terrible experience.",
    "Arrived damaged and customer service was unhelpful. Took three weeks to get a replacement. Very frustrating.",
    "The product is nothing like described. Cheap materials and poor build quality. Would not recommend.",
    "Extremely slow performance and crashes frequently. Complete waste of money. Returning immediately.",
    "Terrible product! The connectivity drops constantly and the software is buggy. Very disappointing purchase.",
    "Poor quality. The charging port broke within a month. Customer support was no help at all. Avoid this product.",
    "Worst purchase ever. Does not match the product description at all. The packaging was deceptive.",
    "Battery life is terrible, lasts only 2 hours. The display is dim and washed out. Very unhappy.",
    "Stopped working after a week. Complete junk. The company refused to honor the warranty. Avoid!",
]

NEUTRAL_REVIEWS = [
    "The product is decent. Does what it says on the box. Nothing special but gets the job done.",
    "Average product. Some features work well, others are mediocre. Might be better options available.",
    "It's okay for the price. Not the best quality but acceptable. Setup took longer than expected.",
    "Works as advertised. No major issues so far. The design could be better but functionality is fine.",
    "Mediocre product. The performance is inconsistent. Sometimes works great, other times struggles.",
    "Three stars. Some good features and some bad ones. The sound quality is okay but not impressive.",
    "Mixed feelings about this purchase. Delivery was fast but the product quality is average at best.",
    "Meets basic requirements. Build quality is average. If you need something basic it will do the job.",
    "Neither good nor bad. The product works but lacks the premium feel the price suggests.",
    "Acceptable for casual use. Heavy users might want something better. Decent starter product.",
]

PRODUCTS = [
    'Wireless Headphones', 'Smart Speaker', 'Laptop', 'Tablet', 'Smartphone',
    'Smartwatch', 'Wireless Earbuds', 'Bluetooth Keyboard', 'USB-C Hub', 'Webcam',
    'Monitor', 'Gaming Mouse', 'Mechanical Keyboard', 'SSD Drive', 'Power Bank',
    'Smart Home Hub', 'Security Camera', 'LED Strip Lights', 'Dash Cam', 'E-Reader'
]

CATEGORIES = {
    'Wireless Headphones': 'Audio', 'Smart Speaker': 'Audio', 'Wireless Earbuds': 'Audio',
    'Laptop': 'Computers', 'Tablet': 'Computers', 'USB-C Hub': 'Computers', 'SSD Drive': 'Computers',
    'Smartphone': 'Mobile', 'Smartwatch': 'Mobile', 'Power Bank': 'Mobile',
    'Bluetooth Keyboard': 'Peripherals', 'Webcam': 'Peripherals', 'Monitor': 'Peripherals',
    'Gaming Mouse': 'Peripherals', 'Mechanical Keyboard': 'Peripherals',
    'Smart Home Hub': 'Smart Home', 'Security Camera': 'Smart Home', 'LED Strip Lights': 'Smart Home',
    'Dash Cam': 'Automotive', 'E-Reader': 'E-Readers'
}

# — Generate 2,000 synthetic reviews —
n = 2000
users = [f'user_{i:04d}' for i in range(1, 401)] # 400 unique users

rows = []
start_ts = 1420070400 # 2015-01-01
end_ts = 1704067200 # 2024-01-01

for i in range(n):
    sentiment_class = random.choices(['positive', 'negative', 'neutral'], weights=[0.55, 0.25, 0.20])[0]
    if sentiment_class == 'positive':
        review_text = random.choice(POSITIVE_REVIEWS)
        rating = random.choices([4, 5], weights=[0.3, 0.7])[0]
    elif sentiment_class == 'negative':
        review_text = random.choice(NEGATIVE_REVIEWS)
        rating = random.choices([1, 2], weights=[0.6, 0.4])[0]
    else:
        review_text = random.choice(NEUTRAL_REVIEWS)
        rating = 3
```

```

product = random.choice(PRODUCTS)
user     = random.choice(users)
ts       = random.randint(start_ts, end_ts)
date     = datetime.utcfromtimestamp(ts)

rows.append({
    'reviewerID' : user,
    'asin'       : f'B{random.randint(1000000,9999999):07d}',
    'productName' : product,
    'category'    : CATEGORIES[product],
    'reviewText'  : review_text,
    'overall'     : rating,
    'unixReviewTime': ts,
    'reviewDate'  : date,
    'year'        : date.year,
    'month'       : date.month,
    'helpful_votes': random.randint(0, 150),
    'total_votes' : random.randint(0, 200),
    'verified'    : random.choice([True, False]),
})

df = pd.DataFrame(rows).sort_values('unixReviewTime').reset_index(drop=True)

print(f" Dataset created: {df.shape[0]:,} rows x {df.shape[1]} columns")
print(f"   Products   : {df['productName'].nunique()}")
print(f"   Users       : {df['reviewerID'].nunique()}")
print(f"   Avg Rating   : {df['overall'].mean():.2f}")
print(f"   Date Range   : {df['reviewDate'].min().date()} → {df['reviewDate'].max().date()}")
display(df.head(3))

```

Dataset created: 2,000 rows x 13 columns 2.0/2.0 MB 11.7 MB/s eta 0:00:00

Products : 20
 Users : 395
 Avg Rating : 3.57
 Date Range : 2015-01-02 → 2023-12-31

	reviewerID	asin	productName	category	reviewText	overall	unixReviewTime	reviewDate	year	month	helpful_votes	total_votes
0	user_0298	B3461642	Smartphone	Mobile	Very disappointed with this purchase. The prod...	1	1420169364	2015-01-02 03:29:24	2015	1	38	

```

for resource in ['punkt', 'punkt_tab', 'stopwords', 'wordnet',
                 'averaged_perceptron_tagger', 'averaged_perceptron_tagger_eng',
                 'vader_lexicon', 'omw-1.4']:
    nltk.download(resource, quiet=True)

```

NLP Preprocessing Pipeline Full pipeline: Tokenization → Lowercasing → Stopword Removal → Lemmatization → POS Tagging

```

def preprocess_text(text: str) -> dict:
    """Full NLP preprocessing pipeline returning multiple representations."""
    # 1. Lowercase
    text_lower = text.lower()
    # 2. Remove punctuation & special chars
    text_clean = re.sub(r'^a-z\s]', '', text_lower)
    # 3. Tokenize
    tokens = word_tokenize(text_clean)
    # 4. Remove stopwords
    tokens_no_stop = [t for t in tokens if t not in STOP_WORDS and len(t) > 2]
    # 5. Lemmatize
    tokens_lemma = [lemmatizer.lemmatize(t) for t in tokens_no_stop]
    # 6. Rejoin
    processed = ' '.join(tokens_lemma)
    return {
        'tokens_raw' : tokens,
        'tokens_clean': tokens_no_stop,
        'tokens_lemma': tokens_lemma,
        'processed' : processed,
        'char_count' : len(text),
        'word_count' : len(tokens),
    }

```

```

        'unique_words': len(set(tokens_lemma)),
    }

print("Running NLP preprocessing pipeline...")
pipeline_results = df['reviewText'].apply(preprocess_text)

df['processed_text'] = pipeline_results.apply(lambda x: x['processed'])
df['word_count']     = pipeline_results.apply(lambda x: x['word_count'])
df['unique_words']   = pipeline_results.apply(lambda x: x['unique_words'])
df['char_count']     = pipeline_results.apply(lambda x: x['char_count'])

# — Sentiment Scoring —
print("Computing VADER sentiment scores...")
df['vader_compound'] = df['reviewText'].apply(lambda x: sia.polarity_scores(x)['compound'])
df['vader_pos']      = df['reviewText'].apply(lambda x: sia.polarity_scores(x)['pos'])
df['vader_neg']      = df['reviewText'].apply(lambda x: sia.polarity_scores(x)['neg'])
df['vader_neu']      = df['reviewText'].apply(lambda x: sia.polarity_scores(x)['neu'])
df['sentiment_label'] = df['vader_compound'].apply(
    lambda x: 'Positive' if x >= 0.05 else ('Negative' if x <= -0.05 else 'Neutral')
)

# TextBlob subjectivity
df['subjectivity'] = df['reviewText'].apply(lambda x: TextBlob(x).sentiment.subjectivity)

print(f"\n Preprocessing complete!")
print(f"  Sentiment distribution:")
print(df['sentiment_label'].value_counts().to_string())
display(df[['reviewText', 'processed_text', 'vader_compound', 'sentiment_label']].head(3))

```

Running NLP preprocessing pipeline...
Computing VADER sentiment scores...

Preprocessing complete!
Sentiment distribution:
sentiment_label
Positive 1372
Negative 598
Neutral 30

	reviewText	processed_text	vader_compound	sentiment_label
0	Very disappointed with this purchase. The prod...	disappointed purchase product stopped working ...	-0.8122	Negative
1	Perfect gift! The recipient loved it. Quality ...	perfect gift recipient loved quality exception...	0.8955	Positive
2	Arrived damaged and customer service was unhel...	arrived damaged customer service unhelpful too...	-0.7264	Negative



Visualization Design

2.1 Exploratory Data Analysis

```

fig, axes = plt.subplots(2, 3, figsize=(18, 10))
fig.suptitle('Amazon Electronics Reviews – Exploratory Data Analysis', fontsize=16, fontweight='bold', y=1.01)

PALETTE = {'Positive': '#2ecc71', 'Neutral': '#f39c12', 'Negative': '#e74c3c'}

# 1. Rating Distribution
ax = axes[0, 0]
rating_counts = df['overall'].value_counts().sort_index()
bars = ax.bar(rating_counts.index, rating_counts.values,
              color=['#e74c3c', '#e67e22', '#f1c40f', '#27ae60', '#2ecc71'], edgecolor='white', linewidth=1.2)
for bar, val in zip(bars, rating_counts.values):
    ax.text(bar.get_x()+bar.get_width()/2, bar.get_height()+10, str(val), ha='center', fontsize=10, fontweight='bold')
ax.set_title('Rating Distribution', fontweight='bold')
ax.set_xlabel('Star Rating'); ax.set_ylabel('Count')
ax.set_xticks([1,2,3,4,5])

# 2. Sentiment Label Distribution
ax = axes[0, 1]
sent_counts = df['sentiment_label'].value_counts()
colors_pie = [PALETTE[s] for s in sent_counts.index]
wedges, texts, autotexts = ax.pie(sent_counts, labels=sent_counts.index, autopct='%1.1f%%',
                                colors=colors_pie, startangle=90,
                                wedgeprops=dict(edgecolor='white', linewidth=2))
for at in autotexts: at.set_fontweight('bold')
ax.set_title('VADER Sentiment Distribution', fontweight='bold')

```

```

# 3. Reviews Over Time
ax = axes[0, 2]
time_series = df.groupby('year').size().reset_index(name='count')
ax.fill_between(time_series['year'], time_series['count'], alpha=0.3, color='#3498db')
ax.plot(time_series['year'], time_series['count'], 'o-', color='#2980b9', linewidth=2.5, markersize=6)
ax.set_title('Reviews Over Time', fontweight='bold')
ax.set_xlabel('Year'); ax.set_ylabel('Number of Reviews')

# 4. Sentiment by Product Category
ax = axes[1, 0]
cat_sent = df.groupby(['category', 'sentiment_label']).size().unstack(fill_value=0)
cat_sent.plot(kind='barh', ax=ax, color=[PALETTE.get(c, 'gray') for c in cat_sent.columns],
              edgecolor='white', linewidth=0.8)
ax.set_title('Sentiment by Category', fontweight='bold')
ax.set_xlabel('Count'); ax.legend(title='Sentiment', fontsize=8)

# 5. VADER Compound Score Distribution
ax = axes[1, 1]
for label, color in PALETTE.items():
    subset = df[df['sentiment_label'] == label]['vader_compound']
    ax.hist(subset, bins=20, alpha=0.6, color=color, label=label, edgecolor='white')
ax.axvline(x=0.05, color='gray', linestyle='--', linewidth=1, alpha=0.7)
ax.axvline(x=-0.05, color='gray', linestyle='--', linewidth=1, alpha=0.7)
ax.set_title('VADER Compound Score', fontweight='bold')
ax.set_xlabel('Compound Score'); ax.set_ylabel('Frequency')
ax.legend()

# 6. Word Count Distribution
ax = axes[1, 2]
ax.hist(df['word_count'], bins=25, color='#9b59b6', alpha=0.8, edgecolor='white')
ax.axvline(df['word_count'].mean(), color='red', linestyle='--', linewidth=2,
           label=f"Mean: {df['word_count'].mean():.0f}")
ax.set_title('Review Word Count', fontweight='bold')
ax.set_xlabel('Word Count'); ax.set_ylabel('Frequency')
ax.legend()

plt.tight_layout()
plt.savefig('eda_overview.png', dpi=150, bbox_inches='tight')
plt.show()
print("EDA plots saved.")

```

Amazon Electronics Reviews — Exploratory Data Analysis



EDA plots saved.

2.2 Word Cloud Visualizations (Baseline + Sentiment-Stratified)

```
def generate_wordcloud(text_series, title, colormap='viridis', background='white',
                      mask_color=None, ax=None):
    """Generate a styled word cloud."""
    combined = ' '.join(text_series.dropna().tolist())
    wc = WordCloud(
        width=800, height=400,
        background_color=background,
        colormap=colormap,
        stopwords=STOPWORDS | STOP_WORDS,
        max_words=200,
        prefer_horizontal=0.8,
        min_font_size=8,
        collocations=True
    ).generate(combined)
    if ax:
        ax.imshow(wc, interpolation='bilinear')
        ax.axis('off')
        ax.set_title(title, fontsize=12, fontweight='bold', pad=10)
    return wc

fig, axes = plt.subplots(2, 2, figsize=(18, 10))
fig.suptitle('Word Cloud Analysis - Amazon Electronics Reviews', fontsize=16, fontweight='bold')

generate_wordcloud(df['processed_text'], 'All Reviews', 'plasma', 'white', ax=axes[0,0])
generate_wordcloud(df['processed_text'], 'Positive', 'plasma', 'white', ax=axes[0,1])
generate_wordcloud(df['processed_text'], 'Negative', 'plasma', 'white', ax=axes[1,0])
generate_wordcloud(df['processed_text'], 'Neutral', 'plasma', 'white', ax=axes[1,1])
```

Word Clouds generated.

```
# — TF-IDF Frequency Plots per Sentiment —
fig, axes = plt.subplots(1, 3, figsize=(20, 7))
fig.suptitle('Top 15 TF-IDF Terms by Sentiment Class', fontsize=15, fontweight='bold')

sentiment_colors = {'Positive': '#27ae60', 'Negative': '#c0392b', 'Neutral': '#e67e22'}

for ax, sentiment in zip(axes, ['Positive', 'Negative', 'Neutral']):
    subset = df[df['sentiment_label'] == sentiment]['processed_text'].fillna('')
    if len(subset) < 2:
        continue
```



```


tfidf = TfidfVectorizer(max_features=15, ngram_range=(1, 2))
tfidf.fit(subset)
scores = dict(zip(tfidf.get_feature_names_out(),
                  tfidf.idf_))

# Use term frequency instead of idf for clearer visualization
cv = CountVectorizer(max_features=15, ngram_range=(1,2))
X = cv.fit_transform(subset)
freq = dict(zip(cv.get_feature_names_out(), X.toarray().sum(axis=0)))
sorted_freq = sorted(freq.items(), key=lambda x: x[1], reverse=True)[:15]
terms, counts = zip(*sorted_freq)

color = sentiment_colors[sentiment]
bars = ax.barh(range(len(terms)), counts, color=color, alpha=0.85, edgecolor='white')
ax.set_yticks(range(len(terms)))
ax.set_yticklabels(terms, fontsize=10)
ax.invert_yaxis()
ax.set_title(f'{sentiment} Reviews', fontweight='bold', color=color, fontsize=13)
ax.set_xlabel('Term Frequency')
for bar, val in zip(bars, counts):
    ax.text(val + 1, bar.get_y() + bar.get_height()/2, str(val),
            va='center', fontsize=8)

plt.tight_layout()
plt.savefig('tfidf_terms.png', dpi=150, bbox_inches='tight')
plt.show()
print(" TF-IDF frequency plots saved.")

```



2.4 N-gram Analysis (Bigrams & Trigrams)

```

from sklearn.feature_extraction.text import CountVectorizer

def get_top_ngrams(corpus, n=2, top_k=15):
    vec = CountVectorizer(ngram_range=(n, n), max_features=top_k*3).fit(corpus)
    bag = vec.transform(corpus)
    sum_words = bag.sum(axis=0)
    words_freq = [(word, sum_words[0, idx]) for word, idx in vec.vocabulary_.items()]
    words_freq = sorted(words_freq, key=lambda x: x[1], reverse=True)
    return words_freq[:top_k]

corpus = df['processed_text'].fillna('').tolist()
bigrams = get_top_ngrams(corpus, n=2)
trigrams = get_top_ngrams(corpus, n=3)

fig, axes = plt.subplots(1, 2, figsize=(18, 7))
fig.suptitle('N-gram Frequency Analysis', fontsize=15, fontweight='bold')


```

```

for ax, ngrams, title, color in [
    (axes[0], bigrams, 'Top 15 Bigrams', '#3498db'),
    (axes[1], trigrams, 'Top 15 Trigrams', '#9b59b6')
]:
    terms, counts = zip(*ngrams)
    bars = ax.barh(range(len(terms)), counts, color=color, alpha=0.85, edgecolor='white')
    ax.set_yticks(range(len(terms)))
    ax.set_yticklabels(terms, fontsize=10)
    ax.invert_yaxis()
    ax.set_title(title, fontweight='bold', fontsize=13)
    ax.set_xlabel('Frequency')
    for bar, val in zip(bars, counts):
        ax.text(val + 0.5, bar.get_y() + bar.get_height()/2, str(int(val)),
            va='center', fontsize=9)

plt.tight_layout()
plt.savefig('ngrams.png', dpi=150, bbox_inches='tight')
plt.show()
print(" N-gram analysis complete.")

```



✓ Topic Modeling (LDA)

Latent Dirichlet Allocation — Topic Extraction

```

# — LDA Topic Modeling —
N_TOPICS = 6

count_vectorizer = CountVectorizer(
    max_df=0.95, min_df=2, max_features=1000, stop_words='english'
)
doc_term_matrix = count_vectorizer.fit_transform(df['processed_text'].fillna(''))
feature_names = count_vectorizer.get_feature_names_out()

lda = LatentDirichletAllocation(
    n_components=N_TOPICS, max_iter=15, learning_method='online',
    learning_offset=50.0, random_state=42
)
lda.fit(doc_term_matrix)

```

```

# Topic labels (human-interpretable)
TOPIC_LABELS = [
    'Sound & Audio Quality',
    'Performance & Speed',
    'Build Quality & Design',
    'Battery & Charging',
    'Customer Service & Delivery',
    'Value for Money'
]

def print_top_words(model, feature_names, n_top_words=10):
    print("\n LDA Topic Keywords:")
    print("-" * 60)
    topics_dict = {}
    for idx, topic in enumerate(model.components_):
        top_words = [feature_names[i] for i in topic.argsort()[::-n_top_words - 1:-1]]
        label = TOPIC_LABELS[idx] if idx < len(TOPIC_LABELS) else f'Topic {idx}'
        print(f" Topic {idx+1} [{label}]: {' '.join(top_words)}")
        topics_dict[idx] = {'label': label, 'words': top_words}
    return topics_dict

topics_dict = print_top_words(lda, feature_names)

# Assign dominant topic to each review
doc_topic_dist = lda.transform(doc_term_matrix)
df['dominant_topic'] = doc_topic_dist.argmax(axis=1)
df['topic_probability'] = doc_topic_dist.max(axis=1)
df['topic_label'] = df['dominant_topic'].map(lambda x: TOPIC_LABELS[x] if x < len(TOPIC_LABELS) else f'Topic {x}')

print(f"\n Topic modeling complete!")
print(f" Topic distribution:")
print(df['topic_label'].value_counts().to_string())

```

LDA Topic Keywords:

```

-----
Topic 1 [Sound & Audio Quality]: purchase, performance, best, setup, price, quality, customer, service, topnotch, year
Topic 2 [Performance & Speed]: quality, time, arrived, recipient, exceptional, loved, wellpackaged, perfect, gift, exactly
Topic 3 [Build Quality & Design]: feature, star, premium, performance, excellent, exceeded, expectation, bad, good, okay
Topic 4 [Battery & Charging]: product, connectivity, great, display, use, fast, sharp, bright, processing, disappoints
Topic 5 [Customer Service & Delivery]: product, work, excellent, purchase, described, delivery, fast, recommend, money, value
Topic 6 [Value for Money]: product, quality, work, battery, life, build, outstanding, amazing, absolutely, perfectly

```

Topic modeling complete!

Topic distribution:

topic_label	
Customer Service & Delivery	461
Sound & Audio Quality	394
Battery & Charging	375
Value for Money	338
Performance & Speed	285
Build Quality & Design	147

3.1 Topic Visualization — Heatmap, Bar Charts, and Bubble Chart

```

# — Topic-Word Heatmap —
N_WORDS = 10
topic_word_matrix = []
all_top_words = set()
for topic in lda.components_:
    top_idx = topic.argsort()[::-N_WORDS-1:-1]
    for i in top_idx:
        all_top_words.add(feature_names[i])

word_list = sorted(list(all_top_words))[:40]
word_idx = {w: i for i, w in enumerate(feature_names)}

heatmap_data = np.zeros((N_TOPICS, len(word_list)))
for t_idx, topic in enumerate(lda.components_):
    for w_idx, word in enumerate(word_list):
        if word in word_idx:
            heatmap_data[t_idx, w_idx] = topic[word_idx[word]]

fig, axes = plt.subplots(1, 2, figsize=(22, 7))
fig.suptitle('LDA Topic Modeling Analysis', fontsize=15, fontweight='bold')

# Heatmap

```


```

sns.heatmap(
    heatmap_data,
    xticklabels=word_list,
    yticklabels=[f'T{i+1}: {TOPIC_LABELS[i][:18]}' for i in range(N_TOPICS)],
    cmap='YlOrRd', ax=axes[0], linewidths=0.1, cbar_kws={'label': 'Word Weight'})
)
axes[0].set_title('Topic-Word Weight Heatmap', fontweight='bold')
axes[0].tick_params(axis='x', rotation=45, labels=7)
axes[0].tick_params(axis='y', labels=9)

# Topic Distribution Bar
topic_counts = df['topic_label'].value_counts()
colors = plt.cm.Set2(np.linspace(0, 1, len(topic_counts)))
bars = axes[1].barh(topic_counts.index, topic_counts.values, color=colors, edgecolor='white', linewidth=1)
axes[1].set_title('Document Distribution Across Topics', fontweight='bold')
axes[1].set_xlabel('Number of Reviews')
for bar, val in zip(bars, topic_counts.values):
    axes[1].text(val + 5, bar.get_y() + bar.get_height()/2, str(val), va='center', fontsize=10)

plt.tight_layout()
plt.savefig('topic_analysis.png', dpi=150, bbox_inches='tight')
plt.show()
print(" Topic visualization saved.")

```



3.2 t-SNE / UMAP Document Embedding Visualization

```

# — t-SNE on TF-IDF features —
tfidf_viz = TfidfVectorizer(max_features=300, stop_words='english')
X_tfidf = tfidf_viz.fit_transform(df['processed_text'].fillna(''))

# Reduce to 50D with SVD first (for speed)
svd = TruncatedSVD(n_components=50, random_state=42)
X_svd = svd.fit_transform(X_tfidf)

print("Running t-SNE (this takes ~30 seconds)...")
tsne = TSNE(n_components=2, perplexity=40, n_iter=600, random_state=42, verbose=0)
X_tsne = tsne.fit_transform(X_svd)

df['tsne_x'] = X_tsne[:, 0]
df['tsne_y'] = X_tsne[:, 1]

fig, axes = plt.subplots(1, 2, figsize=(18, 7))
fig.suptitle('t-SNE Document Embeddings', fontsize=15, fontweight='bold')

# Colored by Sentiment
for label, color in [('Positive', '#27ae60'), ('Neutral', '#f39c12'), ('Negative', '#e74c3c')]:
    mask = df['sentiment_label'] == label
    axes[0].scatter(df.loc[mask, 'tsne_x'], df.loc[mask, 'tsne_y'],

```

```


        c=color, label=label, alpha=0.5, s=15, edgecolors='none')
axes[0].set_title('Colored by Sentiment', fontweight='bold')
axes[0].legend(markerscale=2); axes[0].axis('off')

# Colored by Topic
topic_colors = plt.cm.tab10(np.linspace(0, 1, N_TOPICS))
for idx in range(N_TOPICS):
    mask = df['dominant_topic'] == idx
    axes[1].scatter(df.loc[mask, 'tsne_x'], df.loc[mask, 'tsne_y'],
                    c=[topic_colors[idx]], label=TOPIC_LABELS[idx][:20],
                    alpha=0.5, s=15, edgecolors='none')
axes[1].set_title('Colored by Topic', fontweight='bold')
axes[1].legend(markerscale=2, fontsize=7, loc='lower right'); axes[1].axis('off')

plt.tight_layout()
plt.savefig('tsne_embedding.png', dpi=150, bbox_inches='tight')
plt.show()
print(" t-SNE visualization saved.")

```

Running t-SNE (this takes ~30 seconds)...



Phase 2 (Network): Network Visualization

User-Product Co-Review Graph (Node-Link Diagram)

```

# — Build User-Product Bipartite Graph —
# Filter to most active users for clarity
top_users = df['reviewerID'].value_counts().head(60).index.tolist()
top_products = df['productName'].value_counts().head(20).index.tolist()

df_net = df[df['reviewerID'].isin(top_users) & df['productName'].isin(top_products)].copy()

G = nx.DiGraph()

# Add user nodes
for user in df_net['reviewerID'].unique():
    user_df = df_net[df_net['reviewerID'] == user]
    G.add_node(user, node_type='user',
               avg_sentiment=user_df['vader_compound'].mean(),
               review_count=len(user_df),
               label=user[:8])

```

```

# Add product nodes
for product in df_net['productName'].unique():
    prod_df = df_net[df_net['productName'] == product]
    G.add_node(product, node_type='product',
               avg_rating=prod_df['overall'].mean(),
               review_count=len(prod_df),
               category=CATEGORIES.get(product, 'Other'),
               label=product)

# Add edges (user reviewed product)
for _, row in df_net.iterrows():
    if G.has_edge(row['reviewerID'], row['productName']):
        G[row['reviewerID']][row['productName']]['weight'] += 1
    else:
        G.add_edge(row['reviewerID'], row['productName'],
                   weight=1,
                   sentiment=row['vader_compound'],
                   rating=row['overall'])

# — Compute Graph Metrics —
degree_cent = nx.degree_centrality(G)
betweenness_cent = nx.betweenness_centrality(G, k=min(50, len(G.nodes())))
in_degree = dict(G.in_degree())
out_degree = dict(G.out_degree())

nx.set_node_attributes(G, degree_cent, 'degree_centrality')
nx.set_node_attributes(G, betweenness_cent, 'betweenness_centrality')

print(f" Graph constructed:")
print(f" Nodes: {G.number_of_nodes()} | Edges: {G.number_of_edges()}")
print(f" Density: {nx.density(G):.4f}")
print(f" Top nodes by degree centrality:")
top_nodes = sorted(degree_cent.items(), key=lambda x: x[1], reverse=True)[:5]
for node, score in top_nodes:
    print(f" {node}: {score:.4f}")

```

```

Graph constructed:
Nodes: 80 | Edges: 420
Density: 0.0665
Top nodes by degree centrality:
Smart Speaker: 0.3671
LED Strip Lights: 0.3418
E-Reader: 0.3291
Mechanical Keyboard: 0.3038
Bluetooth Keyboard: 0.2911

```

```

# — Static Network Plot (Matplotlib) – Node-Link Diagram —
fig, ax = plt.subplots(1, 1, figsize=(16, 11))
ax.set_facecolor('#0d1117')
fig.patch.set_facecolor('#0d1117')

pos = nx.spring_layout(G, k=0.8, iterations=40, seed=42)

user_nodes = [n for n, d in G.nodes(data=True) if d.get('node_type') == 'user']
product_nodes = [n for n, d in G.nodes(data=True) if d.get('node_type') == 'product']

# Node sizes proportional to centrality
user_sizes = [300 + degree_cent.get(n, 0) * 2000 for n in user_nodes]
product_sizes = [500 + degree_cent.get(n, 0) * 3000 for n in product_nodes]

# User node colors by avg sentiment
user_sentiment = [G.nodes[n].get('avg_sentiment', 0) for n in user_nodes]

nx.draw_networkx_nodes(G, pos, nodelist=user_nodes, node_size=user_sizes,
                      node_color=user_sentiment, cmap='RdYlGn',
                      vmin=-1, vmax=1, alpha=0.85, ax=ax)
nx.draw_networkx_nodes(G, pos, nodelist=product_nodes, node_size=product_sizes,
                      node_color='#f39c12', alpha=0.9, node_shape='s', ax=ax)
nx.draw_networkx_edges(G, pos, alpha=0.15, edge_color='aaaaaa',
                      width=0.6, arrows=True, arrowsize=6, ax=ax)

# Labels only for products
product_labels = {n: n[:12] for n in product_nodes}
nx.draw_networkx_labels(G, pos, labels=product_labels,
                      font_size=7, font_color='white', ax=ax)


```

```

legend_elements = [
    mpatches.Patch(color='#27ae60', label='User (Positive Sentiment)'),
    mpatches.Patch(color='#e74c3c', label='User (Negative Sentiment)'),
    mpatches.Patch(color='#f39c12', label='Product'),
]
ax.legend(handles=legend_elements, loc='upper left',
          facecolor='#1a1a2e', labelcolor='white', fontsize=9)
ax.set_title('User-Product Review Network (Sentiment-Colored)',
            fontsize=14, color='white', fontweight='bold', pad=15)
ax.axis('off')

plt.tight_layout()
plt.savefig('network_nodelink.png', dpi=150, bbox_inches='tight', facecolor='#0d1117')
plt.show()
print(" Node-link network diagram saved.")

```



Node-link network diagram saved.

Network: Adjacency Matrix (Alternative Representation)

— Adjacency Matrix —


```
# Use only top 15 products and top 20 users for clarity
top15_products = df['productName'].value_counts().head(15).index.tolist()
top20_users    = df['reviewerID'].value_counts().head(20).index.tolist()

pivot = df[df['reviewerID'].isin(top20_users) & df['productName'].isin(top15_products)]\
        .groupby(['reviewerID', 'productName'])['overall'].mean().unstack(fill_value=0)

# Reindex to ensure all products
for p in top15_products:
    if p not in pivot.columns:
        pivot[p] = 0
pivot = pivot[top15_products]

fig, ax = plt.subplots(figsize=(14, 9))
im = sns.heatmap(
    pivot,
    annot=True, fmt='.1f',
    cmap='RdYlGn',
    linewidths=0.4,
    linecolor='#cccccc',
    vmin=0, vmax=5,
    cbar_kws={'label': 'Average Rating (0=no review)'},
    ax=ax
)
ax.set_title('Adjacency Matrix: User x Product Average Ratings',
             fontsize=14, fontweight='bold', pad=15)
ax.set_xlabel('Product', fontsize=11)
ax.set_ylabel('User ID', fontsize=11)
ax.tick_params(axis='x', rotation=35, labelsiz=9)
ax.tick_params(axis='y', rotation=0, labelsiz=8)

plt.tight_layout()
plt.savefig('adjacency_matrix.png', dpi=150, bbox_inches='tight')
plt.show()
print(" Adjacency matrix saved.")
```

Interactive Network with PyVis

```
# — Interactive PyVis Network —
# Smaller subgraph for performance in notebook
top_u = df['reviewerID'].value_counts().head(30).index.tolist()
top_p = df['productName'].value_counts().head(12).index.tolist()
df_small = df[df['reviewerID'].isin(top_u) & df['productName'].isin(top_p)]

G_small = nx.Graph()
for user in df_small['reviewerID'].unique():
    u_df = df_small[df_small['reviewerID']==user]
    sc = u_df['vader_compound'].mean()
    color = '#27ae60' if sc > 0.05 else ('#e74c3c' if sc < -0.05 else '#f39c12')
    G_small.add_node(user, title=f"User: {user}\nReviews: {len(u_df)}\nSentiment: {sc:.2f}",
                     color=color, size=12, group=1)

for prod in df_small['productName'].unique():
```

```

p_df = df_small[df_small['productName']==prod]
avg_r = p_df['overall'].mean()
G_small.add_node(prod, title=f"Product: {prod}\nReviews: {len(p_df)}\nAvg Rating: {avg_r:.1f}",
                  color='#3498db', size=25, group=2)

for _, row in df_small.iterrows():
    G_small.add_edge(row['reviewerID'], row['productName'],
                     title=f"Rating: {row['overall']}\nSentiment: {row['vader_compound']:.2f}")

net = Network(height='550px', width='100%', bgcolor='#1a1a2e', font_color='white',
              notebook=True, directed=False)
net.set_options("""
var options = {
  "nodes": {"borderWidth": 2, "shadow": true},
  "edges": {"color": {"inherit": "from"}, "smooth": false, "width": 0.8},
  "physics": {"forceAtlas2Based": {"gravitationalConstant": -60, "springLength": 100},
              "minVelocity": 0.75, "solver": "forceAtlas2Based"}
}
""")
net.from_nx(G_small)

html_path = 'interactive_network.html'
net.save_graph(html_path)
print(" Interactive network saved to 'interactive_network.html'")
display(HTML(f'<a href="{html_path}" target="_blank"> Open Interactive Network</a>'))

# Display inline in Colab
with open(html_path, 'r') as f:
    html_content = f.read()
display(HTML(f'<div style="height:560px">{html_content}</div>'))

```

Warning: When `cdn_resources` is 'local' jupyter notebook has issues displaying graphics on chrome/safari. Use `cdn_resources='in`
 Interactive network saved to 'interactive_network.html'

[Open Interactive Network](#)



Temporal Evolution Analysis

```

# — Temporal Trend Analysis —
monthly = df.groupby(['year', 'month']).agg(
    review_count = ('reviewText', 'count'),
    avg_sentiment = ('vader_compound', 'mean'),
    avg_rating = ('overall', 'mean'),

```

```

neiprui_votes = ( neiprui_votes , sum )
).reset_index()
monthly['date'] = pd.to_datetime(monthly[['year','month']].assign(day=1))
monthly = monthly.sort_values('date')

fig = make_subplots(
    rows=3, cols=1,
    subplot_titles=[
        'Monthly Review Volume', ' 🟡 Average Sentiment Over Time',
        'Average Rating Over Time'
    ],
    shared_xaxes=True, vertical_spacing=0.08
)

# Review Volume
fig.add_trace(go.Scatter(
    x=monthly['date'], y=monthly['review_count'],
    mode='lines+markers', fill='tozeroy',
    line=dict(color='#3498db', width=2),
    fillcolor='rgba(52,152,219,0.2)',
    name='Review Count'
), row=1, col=1)


# Sentiment
colors_line = [' #27ae60' if v >= 0.05 else '#e74c3c' if v <= -0.05 else '#f39c12'
               for v in monthly['avg_sentiment']]
fig.add_trace(go.Scatter(
    x=monthly['date'], y=monthly['avg_sentiment'],
    mode='lines+markers',
    line=dict(color='#27ae60', width=2),
    marker=dict(color=colors_line, size=6),
    name='Avg Sentiment'
), row=2, col=1)
fig.add_hline(y=0, line_dash='dash', line_color='gray', opacity=0.5, row=2, col=1)

# Rating
fig.add_trace(go.Scatter(
    x=monthly['date'], y=monthly['avg_rating'],
    mode='lines+markers',
    line=dict(color='#e67e22', width=2),
    name='Avg Rating'
), row=3, col=1)
fig.add_hline(y=3, line_dash='dash', line_color='gray', opacity=0.5, row=3, col=1)

fig.update_layout(
    title_text='Temporal Evolution of Amazon Electronics Reviews (2015-2024)',
    title_font_size=16,
    height=700,
    template='plotly_dark',
    showlegend=False
)
fig.write_html('temporal_evolution.html')
fig.show()
print(" Temporal evolution chart saved.")

```

Temporal Evolution of Amazon Electronics Reviews (2015–2024)



Temporal evolution chart saved.

Integration — Sentiment-Aware Topic Visualization

```
# — Sentiment × Topic Integration —
# Bubble Chart: Topic vs Sentiment vs Volume
topic_sent = df.groupby(['topic_label', 'sentiment_label']).agg(
    count      = ('reviewText', 'count'),
    avg_rating  = ('overall', 'mean'),
    avg_compound= ('vader_compound', 'mean')
).reset_index()

fig = px.scatter(
    topic_sent,
    x='avg_rating', y='avg_compound',
    size='count', color='sentiment_label',
    hover_name='topic_label',
    hover_data={'count': True, 'avg_rating': ':.2f', 'avg_compound': ':.2f'},
    color_discrete_map={'Positive': '#27ae60', 'Negative': '#e74c3c', 'Neutral': '#f39c12'},
    size_max=60,
    title='Topic × Sentiment Integration – Bubble Chart<br><sup>Size = number of reviews | X = avg rating | Y = VADER compound score</sup>',
    labels={'avg_rating': 'Average Star Rating', 'avg_compound': 'VADER Compound Score'},
    template='plotly_dark'
)

fig.add_hline(y=0, line_dash='dot', line_color='gray')
fig.add_vline(x=3, line_dash='dot', line_color='gray')
fig.update_traces(marker=dict(opacity=0.8, line=dict(width=1.5, color='white')))
fig.write_html('sentiment_topic_bubble.html')
fig.show()

# — Sankey: Review Flow (Sentiment → Topic → Rating) —
sent_map  = {'Positive':0, 'Neutral':1, 'Negative':2}
topic_map  = {t:i+3 for i,t in enumerate(df['topic_label'].unique())}
rating_map = {r: i+3+len(topic_map) for i,r in enumerate(sorted(df['overall'].unique()))}

links = []
for (sent, topic), grp in df.groupby(['sentiment_label', 'topic_label']):
```


```

links.append({'source': sent_map[sent], 'target': topic_map[topic], 'value': len(grp)})

labels = list(sent_map.keys()) + list(topic_map.keys())
colors_sankey = ([ '#27ae60', '#f39c12', '#e74c3c' ] +
                  [f'hsl({i*60},70%,50%)' for i in range(len(topic_map))])

sankey_fig = go.Figure(go.Sankey(
    node=dict(label=labels, color=colors_sankey, pad=20, thickness=20,
              line=dict(color='white', width=0.5)),
    link=dict(
        source=[l['source'] for l in links],
        target=[l['target'] for l in links],
        value=[l['value'] for l in links],
        color='rgba(180,180,180,0.2)'
    )
))
sankey_fig.update_layout(
    title='Sankey Diagram: Sentiment → Topic Flow',
    height=500, template='plotly_dark'
)
sankey_fig.write_html('sankey_flow.html')
sankey_fig.show()
print(" Integration visualizations saved.")

```



Sankey Diagram: Sentiment → Topic Flow



Integration visualizations saved.

Interactive Dashboard (Plotly Dash)

Full multi-view linked analytics dashboard

```
# — Prepare Dashboard Data —
dash_df = df[['reviewerID','productName','category','overall','sentiment_label',
              'vader_compound','topic_label','year','word_count','subjectivity',
              'tsne_x','tsne_y','processed_text','reviewText']].copy()

# — Build Dash App —
if JUPYTER_DASH:
    app = JupyterDash(__name__, external_stylesheets=[dbc.themes.DARKLY])
else:
    app = dash.Dash(__name__, external_stylesheets=[dbc.themes.DARKLY])

# — KPI Cards —
def kpi_card(title, value, color):
```

```

return dbc.Card([
    dbc.CardBody([
        html.H5(title, className='card-title', style={'color':'#aaa','fontSize':'12px'}),
        html.H3(value, className='card-text', style={'color':color,'fontWeight':'bold'})
    ])
], style={'background':'#1e2130','border':'1px solid {color}','borderRadius':'8px'})

# — Layout —
app.layout = dbc.Container([
    # Header
    dbc.Row([dbc.Col([
        html.H2(" Amazon Electronics Review Analytics Dashboard",
            style={'color':'#fff','fontWeight':'bold','marginBottom':'0'}),
        html.P("Interactive Visual Analytics | COMP-834 | PAF-IAST",
            style={'color':'#aaa','fontSize':'13px'})
    ])], style={'background':'#0d1117','padding':'15px 20px','marginBottom':'10px'}),

    # KPI Row
    dbc.Row([
        dbc.Col(kpi_card('Total Reviews', f"{len(df):,}", '#3498db'), width=2),
        dbc.Col(kpi_card('Unique Users', f"{df['reviewerID'].nunique()}", '#9b59b6'), width=2),
        dbc.Col(kpi_card('Products', f"{df['productName'].nunique()}", '#e67e22'), width=2),
        dbc.Col(kpi_card('Avg Rating', f"{df['overall'].mean():.2f} ", '#f1c40f'), width=2),
        dbc.Col(kpi_card('Positive Reviews', f"{(df['sentiment_label']=='Positive').sum()}", '#27ae60'), width=2),
        dbc.Col(kpi_card('Negative Reviews', f"{(df['sentiment_label']=='Negative').sum()}", '#e74c3c'), width=2),
    ], className='mb-3'),

    # — Sidebar + Main —
    dbc.Row([
        # Sidebar Filters
        dbc.Col([
            dbc.Card([
                dbc.CardHeader(" 🗖️ Filters", style={'color':'#3498db','fontWeight':'bold'}),
                dbc.CardBody([
                    html.Label("Category", style={'color':'#ccc','fontSize':'12px'}),
                    dcc.Dropdown(
                        id='cat-filter',
                        options=[{'label':'All','value':'All'}] +
                            [{'label':c,'value':c} for c in sorted(df['category'].unique())],
                        value='All', clearable=False,
                        style={'background':'#2c2f3a','color':'#000'}
                    ),
                    html.Br(),
                    html.Label("Sentiment", style={'color':'#ccc','fontSize':'12px'}),
                    dcc.Checklist(
                        id='sent-filter',
                        options=[{'label':f" {s}",'value':s} for s in ['Positive','Neutral','Negative']],
                        value=['Positive','Neutral','Negative'],
                        style={'color':'#ddd','fontSize':'13px'},
                        labelStyle={'display':'block','marginBottom':'5px'}
                    ),
                    html.Br(),
                    html.Label("Year Range", style={'color':'#ccc','fontSize':'12px'}),
                    dcc.RangeSlider(
                        id='year-filter',
                        min=df['year'].min(), max=df['year'].max(),
                        step=1, value=[df['year'].min(), df['year'].max()],
                        marks={y:{'label':str(y),'style':{'color':'#aaa','fontSize':'10px'}}
                            for y in range(df['year'].min(), df['year'].max()+1, 2)},
                        tooltip={'placement':'bottom','always_visible':False}
                    ),
                    html.Br(),
                    html.Label("Star Rating", style={'color':'#ccc','fontSize':'12px'}),
                    dcc.RangeSlider(
                        id='rating-filter',
                        min=1, max=5, step=1,
                        value=[1,5],
                        marks={i:{'label':str(i)+'', 'style':{'color':'#f1c40f','fontSize':'10px'}} for i in range(1,6)}
                    ),
                ])
            ], style={'background':'#1e2130','border':'1px solid #333'})
        ], width=2),

        # Main Charts
        dbc.Col([
            # Row 1: Sentiment + Topic
            dbc.Row([
                dbc.Col([
                    dcc.Graph(id='sentiment-pie', style={'height':'200px'})
                ], width=4)
            ])
        ], width=6)
    ])

```

```

        dbc.Col([dcc.Graph(id='sentiment-pie', style={'height': '280px'})], width=4),
        dbc.Col([dcc.Graph(id='topic-bar', style={'height': '280px'})], width=4),
        dbc.Col([dcc.Graph(id='rating-hist', style={'height': '280px'})], width=4),
    ], className='mb-2'),
    # Row 2: Temporal + t-SNE
    dbc.Row([
        dbc.Col([dcc.Graph(id='temporal-chart', style={'height': '280px'})], width=6),
        dbc.Col([dcc.Graph(id='tsne-scatter', style={'height': '280px'})], width=6),
    ], className='mb-2'),
    # Row 3: Heatmap + Details
    dbc.Row([
        dbc.Col([dcc.Graph(id='cat-heatmap', style={'height': '280px'})], width=6),
        dbc.Col([
            dbc.Card([
                dbc.CardHeader(" Selected Review Details", style={'color': '#3498db'}),
                dbc.CardBody([
                    html.Div(id='details-panel', style={'color': '#ddd', 'fontSize': '12px', 'overflowY': 'auto', 'maxHeight':
                        '150px'})
                ])
            ], style={'background': '#1e2130', 'border': '1px solid #333', 'height': '280px'})
        ], width=6),
    ])
], width=10)
])
], fluid=True, style={'background': '#0d1117', 'minHeight': '100vh', 'padding': '10px'})

# — Callbacks —
@app.callback(
    [Output('sentiment-pie', 'figure'), Output('topic-bar', 'figure'),
     Output('rating-hist', 'figure'), Output('temporal-chart', 'figure'),
     Output('tsne-scatter', 'figure'), Output('cat-heatmap', 'figure'),
     Output('details-panel', 'children')],
    [Input('cat-filter', 'value'), Input('sent-filter', 'value'),
     Input('year-filter', 'value'), Input('rating-filter', 'value')]
)
def update_all(category, sentiments, year_range, rating_range):
    # Filter
    mask = (
        (dash_df['year'] >= year_range[0]) & (dash_df['year'] <= year_range[1]) &
        (dash_df['overall'] >= rating_range[0]) & (dash_df['overall'] <= rating_range[1]) &
        (dash_df['sentiment_label'].isin(sentiments))
    )
    if category != 'All':
        mask &= (dash_df['category'] == category)
    dff = dash_df[mask]

    DARK = {'paper_bgcolor': '#1e2130', 'plot_bgcolor': '#1e2130',
            'font': {'color': '#ccc', 'size': 11}, 'margin': {'l': 40, 'r': 20, 't': 35, 'b': 30}}
    SENT_COLOR = {'Positive': '#27ae60', 'Neutral': '#f39c12', 'Negative': '#e74c3c'}

    # 1. Sentiment Pie
    sc = dff['sentiment_label'].value_counts()
    pie = go.Figure(go.Pie(
        labels=sc.index, values=sc.values,
        marker_colors=[SENT_COLOR.get(l, 'gray') for l in sc.index],
        hole=0.4, textinfo='percent+label'
    ))
    pie.update_layout(title='Sentiment Split', **DARK, showlegend=False)

    # 2. Topic Bar
    tc = dff['topic_label'].value_counts().head(6)
    bar = go.Figure(go.Bar(
        x=tc.index, y=[t[1] for t in tc.index],
        orientation='h',
        marker_color='#3498db',
        marker_line_width=0
    ))
    bar.update_layout(title='Top Topics', **DARK)

    # 3. Rating Histogram
    hist = px.histogram(dff, x='overall', color='sentiment_label',
                        color_discrete_map=SENT_COLOR, nbins=5, barmode='stack')
    hist.update_layout(title='Rating Distribution', **DARK, showlegend=False)

    # 4. Temporal
    tmp = dff.groupby('year').agg(count=('reviewText', 'count'),
                                avg_sent=('vader_compound', 'mean')).reset_index()
    temp_fig = make_subplots(specs=[[{'secondary_y': True}]])
    temp_fig.add_trace(go.Bar(x=tmp['year'], y=tmp['count'], name='Reviews',

```



```

        marker_color='rgba(52,152,219,0.5)'), secondary_y=False)
temp_fig.add_trace(go.Scatter(x=tmp['year'], y=tmp['avg_sent'], name='Sentiment',
                             line=dict(color='#27ae60', width=2), mode='lines+markers'),
                  secondary_y=True)
temp_fig.update_layout(title='Reviews & Sentiment Over Time', **DARK, showlegend=False)

# 5. t-SNE
tsne_fig = px.scatter(
    dff.sample(min(500, len(dff)), random_state=42) if len(dff) > 500 else dff,
    x='tsne_x', y='tsne_y', color='sentiment_label',
    color_discrete_map=SENT_COLOR,
    hover_data={'productName':True, 'overall':True},
    opacity=0.65
)
tsne_fig.update_traces(marker=dict(size=4))
tsne_fig.update_layout(title='t-SNE Document Map', **DARK, showlegend=False)

# 6. Category Heatmap
heat_data = dff.groupby(['category', 'sentiment_label']).size().unstack(fill_value=0)
for s in ['Positive', 'Neutral', 'Negative']:
    if s not in heat_data.columns: heat_data[s] = 0
heat_fig = go.Figure(go.Heatmap(
    z=heat_data.values,
    x=heat_data.columns.tolist(),
    y=heat_data.index.tolist(),
    colorscale='RdYlGn',
    text=heat_data.values, texttemplate='%{text}',
    showscale=True
))
heat_fig.update_layout(title='Category x Sentiment', **DARK)

# 7. Details panel
sample = dff.sample(min(5, len(dff)), random_state=1) if len(dff) > 0 else dff
details = [
    html.Div([
        html.Strong(f"{row['productName']} ({row['overall']})",
                    style={'color': SENT_COLOR.get(row['sentiment_label'], '#fff')}),
        html.P(row['reviewText'][:120] + '...', style={'margin': '2px 0 8px 0', 'color': '#aaa'})
    ]) for _, row in sample.iterrows()
] if len(sample) else [html.P("No data matches current filters.")]

return pie, bar, hist, temp_fig, tsne_fig, heat_fig, details

print(" Dash app defined. Run the cell below to launch it.")

```

Start coding or [generate](#) with AI.

✓ Evaluation

6.A — Quantitative Graph Metrics

```

# — Graph Metrics —
print("=" * 55)
print(" QUANTITATIVE GRAPH EVALUATION METRICS")
print("=" * 55)

UG = G.to_undirected()

# Community detection (modularity)
from networkx.algorithms import community
communities = community.greedy_modularity_communities(UG)
modularity = community.modularity(UG, communities)

# Centrality
deg_cent = nx.degree_centrality(UG)
close_cent = nx.closeness_centrality(UG)
between_cent = nx.betweenness_centrality(UG, k=min(50, len(UG.nodes())))

# Degree distribution
degrees = [d for _, d in UG.degree()]

metrics = {

```

```

'Number of Nodes'      : UG.number_of_nodes(),
'Number of Edges'      : UG.number_of_edges(),
'Graph Density'        : f"{nx.density(UG):.4f}",
'Modularity (communities)' : f"{modularity:.4f}",
'Number of Communities' : len(communities),
'Avg Degree'           : f"{np.mean(degrees):.2f}",
'Max Degree'           : max(degrees),
'Avg Degree Centrality' : f"{np.mean(list(deg_cent.values())):.4f}",
'Avg Betweenness Centrality' : f"{np.mean(list(between_cent.values())):.4f}",
'Avg Closeness Centrality' : f"{np.mean(list(close_cent.values())):.4f}",
'Connected Components' : nx.number_connected_components(UG),
}

for k, v in metrics.items():
    print(f" {k:<35} : {v}")

print()
print("=" * 55)
print(" TEXT ANALYTICS EVALUATION METRICS")
print("=" * 55)

# Topic coherence proxy (avg pairwise PMI via co-occurrence)
text_metrics = {
    'Total Reviews'      : len(df),
    'Avg Word Count'     : f"{df['word_count'].mean():.1f}",
    'Vocabulary Size'    : len(feature_names),
    'Avg VADER Compound' : f"{df['vader_compound'].mean():.3f}",
    'Avg Subjectivity'   : f"{df['subjectivity'].mean():.3f}",
    'Positive Review %'  : f"{(df['sentiment_label']=='Positive').mean()*100:.1f}%",
    'Negative Review %'  : f"{(df['sentiment_label']=='Negative').mean()*100:.1f}%",
    'LDA Log-Likelihood' : f"{lda.score(doc_term_matrix):.2f}",
    'LDA Perplexity'     : f"{lda.perplexity(doc_term_matrix):.2f}",
    'Number of Topics'   : N_TOPICS,
    'Avg Topic Probability' : f"{df['topic_probability'].mean():.3f}",
}

for k, v in text_metrics.items():
    print(f" {k:<35} : {v}")

```

```

=====
QUANTITATIVE GRAPH EVALUATION METRICS
=====
Number of Nodes      : 80
Number of Edges      : 420
Graph Density        : 0.1329
Modularity (communities) : 0.2077
Number of Communities : 4
Avg Degree           : 10.50
Max Degree           : 29
Avg Degree Centrality : 0.1329
Avg Betweenness Centrality : 0.0146
Avg Closeness Centrality : 0.4662
Connected Components : 1

=====
TEXT ANALYTICS EVALUATION METRICS
=====
Total Reviews      : 2000
Avg Word Count     : 16.6
Vocabulary Size    : 171
Avg VADER Compound : 0.343
Avg Subjectivity   : 0.693
Positive Review %  : 68.6%
Negative Review %  : 29.9%
LDA Log-Likelihood : -77458.32
LDA Perplexity     : 55.87
Number of Topics   : 6
Avg Topic Probability : 0.907

```

6.B — Degree Distribution Plot & Network Analysis

```

fig, axes = plt.subplots(1, 3, figsize=(18, 5))
fig.suptitle('Network Evaluation – Structural Analysis', fontsize=14, fontweight='bold')

# 1. Degree Distribution
degree_seq = sorted(degrees, reverse=True)
axes[0].hist(degree_seq, bins=20, color='#3498db', alpha=0.85, edgecolor='white')
axes[0].set_title('Degree Distribution', fontweight='bold')

```

```

axes[0].set_xlabel('Degree'); axes[0].set_ylabel('Count')
axes[0].axvline(np.mean(degree_seq), color='red', linestyle='--',
               label=f'Mean: {np.mean(degree_seq):.1f}')
axes[0].legend()


# 2. Centrality Comparison
top10_nodes = sorted(deg_cent.items(), key=lambda x: x[1], reverse=True)[:10]
nodes_10, deg_10 = zip(*top10_nodes)
bet_10 = [between_cent.get(n, 0) for n in nodes_10]
close_10 = [close_cent.get(n, 0) for n in nodes_10]

x_pos = np.arange(len(nodes_10))
w = 0.25
axes[1].bar(x_pos - w, deg_10, width=w, label='Degree', color='#3498db', alpha=0.85)
axes[1].bar(x_pos, bet_10, width=w, label='Betweenness', color='#e74c3c', alpha=0.85)
axes[1].bar(x_pos + w, close_10, width=w, label='Closeness', color='#27ae60', alpha=0.85)
axes[1].set_xticks(x_pos)
axes[1].set_xticklabels([n[:8] for n in nodes_10], rotation=40, fontsize=8)
axes[1].set_title('Centrality Comparison (Top 10 Nodes)', fontweight='bold')
axes[1].set_ylabel('Centrality Score')
axes[1].legend(fontsize=8)

# 3. Community Sizes
comm_sizes = sorted([len(c) for c in communities], reverse=True)
axes[2].bar(range(len(comm_sizes)), comm_sizes,
            color=plt.cm.tab20(np.linspace(0,1,len(comm_sizes))),
            edgecolor='white')
axes[2].set_title(f'Community Size Distribution\n(Modularity = {modularity:.3f})', fontweight='bold')
axes[2].set_xlabel('Community Index'); axes[2].set_ylabel('Nodes in Community')

plt.tight_layout()
plt.savefig('network_evaluation.png', dpi=150, bbox_inches='tight')
plt.show()
print(" Network evaluation plots saved.")

```



6.C — Simulated User Study (SUS & Task Evaluation)

```

# — Simulated SUS (System Usability Scale) Evaluation —
# In a real study, 5-10 users would complete tasks and answer SUS questionnaire
# Here we simulate plausible results for demonstration

np.random.seed(42)
n_users = 10
SUS_QUESTIONS = [
    'I would use this system frequently',
    'System is unnecessarily complex',
    'System was easy to use',
    'Need technical support to use',
    'Functions were well integrated',
    'Too much inconsistency',
    'Most people would learn quickly',

```

```

    'System was cumbersome to use',
    'Felt very confident using it',
    'Needed to learn many things first'
]

# Odd questions: higher is better (1-5 scale)
# Even questions: lower is better (1-5 scale)
raw_scores = np.zeros((n_users, 10))
for i in range(n_users):
    for q in range(10):
        if q % 2 == 0: # Odd (0-indexed even) - positive
            raw_scores[i, q] = np.clip(np.random.normal(4.1, 0.6), 1, 5)
        else: # Even (0-indexed odd) - negative
            raw_scores[i, q] = np.clip(np.random.normal(1.8, 0.7), 1, 5)

raw_scores = np.round(raw_scores).astype(int)

# SUS Score calculation
def calc_sus(row):
    total = 0
    for i, score in enumerate(row):
        if i % 2 == 0: # odd questions (1,3,5,7,9)
            total += score - 1
        else: # even questions (2,4,6,8,10)
            total += 5 - score
    return total * 2.5

sus_scores = [calc_sus(raw_scores[i]) for i in range(n_users)]
avg_sus = np.mean(sus_scores)

# Task completion times (seconds) - static vs interactive
tasks = ['Find top product', 'Identify negative topics', 'Track yearly trend', 'Compare categories', 'Find influential user']
static_times = np.random.normal([35, 52, 48, 62, 55], 8, (n_users, 5))
interactive_times = np.random.normal([18, 22, 15, 20, 19], 5, (n_users, 5))

static_times = np.clip(static_times, 5, 120)
interactive_times = np.clip(interactive_times, 5, 120)

fig, axes = plt.subplots(1, 3, figsize=(18, 6))
fig.suptitle('User Study Evaluation Results', fontsize=14, fontweight='bold')

# SUS Scores
colors_sus = ['#27ae60' if s >= 70 else '#e67e22' if s >= 50 else '#e74c3c' for s in sus_scores]
axes[0].barh(range(n_users), sus_scores, color=colors_sus, alpha=0.85, edgecolor='white')
axes[0].axvline(avg_sus, color='blue', linestyle='--', linewidth=2, label=f'Avg SUS: {avg_sus:.1f}')
axes[0].axvline(70, color='green', linestyle=':', linewidth=1.5, label='Good (70)')
axes[0].set_yticks(range(n_users)); axes[0].set_yticklabels([f'User {i+1}' for i in range(n_users)])
axes[0].set_title('SUS Scores per User', fontweight='bold')
axes[0].set_xlabel('SUS Score (0-100)'); axes[0].set_xlim(0, 100)
axes[0].legend(fontsize=9)

# Task Completion Time
x = np.arange(len(tasks))
w = 0.35
axes[1].bar(x-w/2, static_times.mean(axis=0), width=w, label='Static', color='#e74c3c', alpha=0.85,
            yerr=static_times.std(axis=0), capsize=5)
axes[1].bar(x+w/2, interactive_times.mean(axis=0), width=w, label='Interactive', color='#27ae60', alpha=0.85,
            yerr=interactive_times.std(axis=0), capsize=5)
axes[1].set_xticks(x); axes[1].set_xticklabels(tasks, rotation=25, ha='right', fontsize=9)
axes[1].set_title('Task Completion Time: Static vs Interactive', fontweight='bold')
axes[1].set_ylabel('Time (seconds)'); axes[1].legend()


# SUS Interpretation
sus_df = pd.DataFrame({'User': [f'U{i+1}' for i in range(n_users)], 'SUS': sus_scores})
bins = [0, 50, 60, 70, 80, 90, 100]
labels_bin = ['Not Acceptable', 'OK', 'Good', 'Good', 'Excellent', 'Excellent']
hist, edges = np.histogram(sus_scores, bins=bins)
axes[2].bar(range(len(hist)), hist, color=['#e74c3c', '#e67e22', '#f1c40f', '#27ae60', '#1abc9c'],
            alpha=0.85, edgecolor='white')
axes[2].set_xticks(range(len(hist)))
axes[2].set_xticklabels([f'{edges[i]:.0f}-{edges[i+1]:.0f}' for i in range(len(hist))], rotation=25)
axes[2].set_title('SUS Score Distribution\nAvg = {avg_sus:.1f} ({\"Excellent\" if avg_sus>=80 else \"Good\" if avg_sus>=70 else \"Not Acceptable\" if avg_sus<70})',
                fontweight='bold')
axes[2].set_ylabel('Number of Users')

plt.tight_layout()
plt.savefig('user_study.png', dpi=150, bbox_inches='tight')

```

```
plt.show()
```

```
print(f"\n{'='*55}")
print(f" USER STUDY SUMMARY")
print(f"{'='*55}")
print(f"   Average SUS Score           : {avg_sus:.1f}/100")
print(f"   SUS Rating                    : {'Excellent' if avg_sus>=80 else 'Good' if avg_sus>=70 else 'OK'}")
print(f"   Avg Static Task Time          : {static_times.mean():.1f}s")
print(f"   Avg Interactive Task Time     : {interactive_times.mean():.1f}s")
```



```
=====
USER STUDY SUMMARY
=====
Average SUS Score           : 77.5/100
SUS Rating                  : Good
Avg Static Task Time        : 50.1s
Avg Interactive Task Time   : 19.2s
Speed Improvement           : 62%
```

✓ Subjectivity-Sentiment Scatter + Treemap

```
# — Advanced Plotly Visualizations —
```

```
# 1. Subjectivity vs Sentiment Scatter (Interactive)
```

```
fig1 = px.scatter(
    df.sample(500, random_state=42),
    x='subjectivity', y='vader_compound',
    color='sentiment_label',
    size='word_count',
    hover_data=['productName', 'overall', 'topic_label'],
    color_discrete_map={'Positive': '#27ae60', 'Neutral': '#f39c12', 'Negative': '#e74c3c'},
    title='Subjectivity vs. Sentiment Compound Score<br><sup>Size = word count | Color = sentiment class</sup>',
    labels={'subjectivity': 'TextBlob Subjectivity', 'vader_compound': 'VADER Compound Score'},
    opacity=0.75, template='plotly_dark'
)
fig1.add_hline(y=0, line_dash='dash', line_color='gray', opacity=0.5)
fig1.add_vline(x=0.5, line_dash='dash', line_color='gray', opacity=0.5)
fig1.write_html('subjectivity_scatter.html')
fig1.show()
```

```
# 2. Treemap: Category → Product → Sentiment
```

```
tree_df = df.groupby(['category', 'productName', 'sentiment_label']).size().reset_index(name='count')
fig2 = px.treemap(
    tree_df,
```